

# Leader Election in rings

All nodes are equal, but some are more equal than others

Marco Aiello, Eirini Kaldeli

University of Groningen

Distributed Systems, 2009

- 1 Then leader election problem
- 2 Ring networks
  - Anonymous vs. non-anonymous rings
- 3 An  $O(n^2)$  LE algorithm for unidirectional rings:  
LeLann-Chang-Roberts
  - Complexity of the LCR LE algorithm
- 4 An  $O(n \log n)$  LE algorithm: Hirschberg-Sinclair
  - Complexity of the HS algorithm


- Want to designate a **unique** processor as a leader, i.e. the coordinator of a task.
- The network nodes communicate in order to make a decision according to some common criterion that breaks the symmetry among them.
- Helpful in achieving fault-tolerance and saving resources. E.g. generate a new single token when a loss is detected in a token ring, or break a deadlock by removing a node from the cycle.
- There are plenty of algorithms appropriate for different network graphs, such as bi-/unidirectional rings, complete graphs, grids etc.
- E.g. given a spanning tree, leader election can be achieved by applying convergecast on it.

- Any process can initiate the LE algorithm (several elections can be called concurrently).
- Every  $p_i$  has two boolean variables *done* and *is\_leader*. *done* is set when  $p_i$  knows that the algorithm has finished, while *is\_leader* is set when  $p_i$  knows that it is the leader.
- An LE algorithm has to satisfy the following two properties:
  - *Safety*: At most one  $p_i$  is a leader:  
$$\forall i, j \ i \neq j : \neg (is\_leader(i) \wedge is\_leader(j))$$
  - *Liveness*: Eventually all  $p_i$ s are either leaders or not and at least one  $p_i$  is a leader:  $\forall i \ done(i) \wedge \exists j \ is\_leader(j)$

# The ring topology

- We will consider a network of  $n$  processors circularly placed on a ring.
- Unidirectional (clockwise): each  $p_i$  sends messages to  $p_{i+1}$  and receives messages from  $p_{i-1}$  (we assume *modulo*  $n$  arithmetics).
- Bidirectional: each  $p_i$  can send and receive messages in both directions.
- Lower bounds and impossibility results for rings also apply to arbitrary topologies.

A clockwise ring



## Leader election in anonymous rings


A ring is **anonymous** if the  $p_i$ s are indistinguishable; they have no unique identifiers, and they all have identical state machines, with the same initial state.

### Theorem

*There is no deterministic leader election algorithm (even) for synchronous anonymous rings (and even for uniform ones).*

### Proof sketch

- All  $p_i$ s start at the same initial state with the same outgoing messages.
- In every round each  $p_i$  sends the same messages to its neighbour, and thus all  $p_i$ s receive exactly the same messages.
- Thus, because all  $p_i$ s have the same state machine, they move to the same state.

- Impossibility of leader election for asynchronous anonymous rings follows.
- Have to introduce some initial asymmetry in the network  processors are assigned identifiers.
- Identifiers have to be unique and totally ordered. Each  $p_i$  knows only its own identifier.
- The algorithms that we will present suit for both synchronous and asynchronous rings.
- We will consider the asynchronous case for our analysis: assume reliable FIFO channels.
- The size of the ring  $n$  is not a priori known to the nodes: **non-uniform** rings.
- At the end of the algorithm the  $p_i$  with the maximal  $id$  is elected, while all  $p_i$ s must know the  $id$  of the elected leader.

- Assume clockwise unidirectional ring.
- One or more  $p_i$ s can take the initiative and start an election, by sending an election message containing their  $id$  to  $p_{i+1}$ .
- When a  $p_i$  spontaneously or upon receiving a message goes in an election, it marks itself as a participant.
- If the  $p_i$  receiving an election message has a greater  $id$  and is not already a participant, then it sends an election message with its own  $id$  to  $p_{i+1}$ .
- If its own  $id$  is smaller, it forwards the message with the  $id$  it has received.
- If it receives a message with its own  $id$  then it declares itself as the leader.



## The LCR algorithm: code for $p_i$ , $0 \leq i \leq n$

```
boolean participant=false;  
int leader_id=null;
```

To initiate an election:

```
send(ELECTION $\langle$ my_id $\rangle$ );  
participant:=true;
```

Upon receiving a message **ELECTION** $\langle j \rangle$ :

```
if ( $j > \textit{my\_id}$ ) then send(ELECTION $\langle j \rangle$ );  
if ( $\textit{my\_id} = j$ ) then send(LEADER $\langle \textit{my\_id} \rangle$ );  
if ( $(\textit{my\_id} > j) \wedge (\neg \textit{participant})$ ) then  
    send(ELECTION $\langle \textit{my\_id} \rangle$ );
```

```
participant:=true;
```


Upon receiving a message **LEADER** $\langle j \rangle$ :

```
leader_id:= $j$ ;  
if ( $\textit{my\_id} \neq j$ ) then send(LEADER $\langle j \rangle$ );
```

- Only the message with the largest identity completes the round trip and returns to its originator, which becomes the leader.
- Time complexity:  $O(n)$
- The leader has to announce itself to all  $p_i$ s through the leader messages, so that termination is guaranteed and everybody knows who the leader is.
- The algorithm verifies the safety and liveness conditions with:
  - $done(i) \equiv (leader\_id(i) \neq \text{null})$
  - $is\_leader(i) \equiv (leader\_id(i) = i)$

## An example run of the LCR algorithm


Assume all  $p_i$ s are initiators.



Messages transmitted:

$\langle 1 \rangle$	1 times
$\langle 2 \rangle$	1 times
$\langle 3 \rangle$	1 times
$\langle 4 \rangle$	1 times
$\langle 5 \rangle$	1 times
<b>total</b>	<b>5 times</b>


## An example run of the LCR algorithm



Messages transmitted:

$\langle 1 \rangle$	1 times
$\langle 2 \rangle$	2 times
$\langle 3 \rangle$	2 times
$\langle 4 \rangle$	2 times
$\langle 5 \rangle$	2 times
<b>total</b>	<b>9 times</b>


# An example run of the LCR algorithm



Messages transmitted:

$\langle 1 \rangle$	1 times
$\langle 2 \rangle$	2 times
$\langle 3 \rangle$	3 times
$\langle 4 \rangle$	3 times
$\langle 5 \rangle$	3 times
<b>total</b>	12 times


# An example run of the LCR algorithm



Messages transmitted:

$\langle 1 \rangle$	1 times
$\langle 2 \rangle$	2 times
$\langle 3 \rangle$	3 times
$\langle 4 \rangle$	4 times
$\langle 5 \rangle$	4 times
<b>total</b>	14 times

## An example run of the LCR algorithm



Messages transmitted:

$\langle 1 \rangle$	1 times
$\langle 2 \rangle$	2 times
$\langle 3 \rangle$	3 times
$\langle 4 \rangle$	4 times
$\langle 5 \rangle$	5 times
<b>total</b>	15 times

Now, the leader  $id = \langle 5 \rangle$  has to be announced to all nodes with 5 more messages. So, in total  $15+5=20$  messages are transmitted.

☞ Note that each identifier  $i$  is sent  $i$  times.

- We are interested in message complexity: Depends on how the ids are arranged.
  - The largest *id* always travels all around the ring ( $n$  msgs).
  - 2nd largest *id* travels until reaching the largest.
  - 3rd largest *id* travels until reaching largest or second largest.
  - ... etc.
- **Worst** way to arrange the ids is in decreasing order (and all  $p_i$ s are initiators): 2nd largest causes  $n - 1$  messages, 3rd largest  $n - 2$  messages etc.
- Number of msgs =  $(n + (n - 1) + \dots + 1) + n = \frac{n(n+1)}{2} + n$   
(including the  $n$  leader messages at the end).
- Worst case complexity =  $O(n^2)$



### Theorem

*The average message complexity of the LCR algorithm is  $O(n \log n)$ .*

### Proof.

- Consider all  $n!$  rings (all possible permutations).
- Each  $id$  makes 1 step  $\rightarrow n!$  times.
- Each  $id$  takes a  $k$ th step if it is the largest among all its neighbours from  $p_{i+1}$  to  $p_{i+k-1}$ :  $\Pr\{\max\_among\_k\} = \frac{1}{k}$ .
- Add  $n!n$  times for the leader announcement phase.
- So, average number of messages =  $\frac{1}{n!}n((n! + \frac{n!}{2} + \dots + \frac{n!}{n}) + n!) = n(1 + \frac{1}{2} + \dots + \frac{1}{n}) + n = O(n)O(\log n) = O(n \log n)$ . □

- Can we improve message complexity?
- There are several algorithms that solve the problem of leader election in asynchronous rings with  $O(n \log n)$  message complexity.
- Try to have messages containing smaller *ids* travel smaller distances across the ring.
- Hirschberg and Sinclair (HS) algorithm: carry out elections on increasingly larger sets of  $p_i$ s.
- Assume that links allow *bidirectional* communication, again  $n$  is not known in advance.

## The HS algorithm: elections in neighbourhoods

- Elections are performed in neighbourhoods: the  **$k$ -neighbourhood** of a  $p_r$  is the set of processors that are at distance at most  $k$  from  $p_r$  ( $k$  left plus  $k$  right neighbours).
- Operate in (asynchronous) phases:  $p_i$  tries to become a leader in phase  $k$  among its  $2^k$  neighbourhood; only if  $p_i$  is the winner, i.e. it has the highest *id* in its  $2^k$  – neighbourhood, it can proceed to phase  $k + 1$ .
- The size of the neighbourhood doubles in each phase.
- Fewer  $p_i$ s proceed to higher phases, until a single winner gets elected in the whole ring.

- Initially, all  $p_i$ s initiate a candidacy (phase 0), e.g. after having received a broadcasted request for electing a leader.
- The ELECTION messages sent by candidates contain three fields:
  - The  $id$  of the candidate.
  - The current phase number  $k$ .
  - A hop counter  $d$ , which is initially 0 and is incremented by 1 whenever the message is forwarded to the next  $p_i$ .
- If a  $p_j$  receives a **ELECTION** $\langle r, k, d \rangle$  where  $d = 2^k$  then it is the last processor in the  $2^k$ -neighbourhood of  $p_r$  with  $id = r$ .

## The HS algorithm: sending messages

- If the  $p_i$  receiving the election message has a greater  $id$ , then it swallows the message, otherwise it relays it to the same direction, after incrementing  $d$  by 1.
- If the message makes it till the  $2^k$ -distance  $p_i$ , then  $p_i$  sends back a REPLY message, which is forwarded till it reaches the candidate  $p_r$ .
- If the candidate receives replies from both directions, then it is the winner of its  $2^k$  neighbourhood.
- A  $p_i$  that receives an election message with its own  $id$  is the leader of the ring.
- The leader should also announce itself to all other nodes (like in LHR).

## The HS algorithm: code for $p_i$ , $1 \leq i \leq n$

To initiate an election (phase 0):

```
send(ELECTION⟨my_id, 0, 0⟩) to left and right;
```

Upon receiving a message ELECTION⟨ $j, k, d$ ⟩ from left (right):

```
if (( $j > my\_id$ )  $\wedge$  ( $d \leq 2^k$ )) then
```

```
    send(ELECTION⟨ $j, k, d + 1$ ⟩) to right (left);
```

```
if (( $j > my\_id$ )  $\wedge$  ( $d = 2^k$ )) then
```

```
    send(REPLY⟨ $j, k$ ⟩) to left (right);
```

```
if ( $my\_id = j$ ) then announce itself as leader;
```

Upon receiving a message REPLY⟨ $j, k$ ⟩ from left (right):

```
if ( $my\_id \neq j$ ) then
```

```
    send(REPLY⟨ $j, k$ ⟩) to right (left);
```

```
else
```

```
    if (already received REPLY⟨ $j, k$ ⟩)
```

```
        send(ELECTION⟨ $j, k + 1, 1$ ⟩) to left and right;
```

- At phase  $k$  at most  $4 \cdot 2^k$  messages are circulated on behalf of a particular candidate (elections and replies).
- How many candidates compete in phase  $k$ , in worst case?
- At phase  $k = 0$  there are  $n$  candidates.

### Lemma

*For every  $k \geq 1$  the number of processors that will continue to phase  $k$  is at most  $\lfloor \frac{n}{2^{k-1}+1} \rfloor$ .*

- Proof: the minimum distance between two winners at phase  $k - 1$  is  $2^{k-1} + 1$ .
- The total number of messages sent at phase  $k$  that is not the last phase is  $4(2^k \lfloor \frac{n}{2^{k-1}+1} \rfloor) = 8n \lfloor \frac{2^{k-1}}{2^{k-1}+1} \rfloor < 8n$

- The total number of phases till the leader is elected is  $\lceil \log n \rceil + 1$  (including phase 0).
- In last phase  $2n$  msgs are sent (no replies).
- So, the total number of messages in worst case is
$$4n + \sum_{k=1}^{\lceil \log n \rceil - 1} (4 \cdot 2^k \frac{n}{2^{k-1} + 1}) + 2n \leq 6n + 8n(\lceil \log n \rceil - 1).$$
- Message complexity:  $O(n \log n)$



- The max time for each phase  $k$  that is not the final is  $2^k$ .
- The max total time required by phases 0 to  $k$  is  $2(2^0 + 2^1 + \dots 2^k) = 2(2^{k+1} - 1)$ .
- The max total time required by all phases till the penultimate one is thus  $2(2^{\lceil \log n \rceil} - 1)$ .
- Time for the final phase is  $n$ .
- Time complexity:  $O(n)$

But, can we do better than  $O(n \log n)$ ?

### Theorem

*Any leader election algorithm for asynchronous rings whose size is not known a priori has  $\Omega(n \log n)$  message complexity (holds also for unidirectional rings).*

- Both LHR and HS are *comparison-based* algorithms, i.e. they use the identifiers only for comparisons ( $<$ ,  $>$ ,  $=$ ).
- In synchronous networks,  $O(n)$  message complexity can be achieved if general arithmetic operations are permitted (non-comparison based) and if time complexity is unbounded.